Problems of statistical power in infancy research

Michael C. Frank

Department of Psychology, Stanford University

Thanks to ...

Please address correspondence to Michael C. Frank, Department of Psychology, Stanford University, 450 Serra Mall (Jordan Hall), Stanford, CA, 94305, tel: (650) 724-4003, email: mcfrank@stanford.edu.

Abstract

Reproducibility is a core value for empirical research. There is increasing concern throughout psychology that a variety of statistical standards and methodological practices have led to low empirical rates of replication. Because research with infants is labor-intensive and slow, reproducibility rates have not been assessed directly. Yet the small sample sizes and limited numbers of measurements available in infant research lead to significant concern about this field. I show here that most infancy research is underpowered, that there is significant publication bias in this field, and that the conventional sample sizes are insufficient for detecting the majority of effects, especially when negative control groups and post-hoc analyses are employed. I end by considering the implications of these findings for researchers and for the literature.

Introduction

Replicability is a cornerstone of empirical science. If an experiment cannot be repeated, it is essentially an anecdote.

The issues with questionable research practices, statistical power, and null-hypothesis statistical testing more broadly are well-known at this point, and my goal here is not to reiterate these discussions. Interested readers are directed to ? (?), ? (?) etc. for tutorial discussion. In the current article, my goal is to focus specifically on the infancy literature—even more specifically on looking time studies—and highlight a number of methodological practices that may lead to a particular inflation of false positives in this literature.

I should note that although I recognize the issues with the null-hypothesis significance testing framwork, I believe it is

Issues specific to infancy research

1. Sample sizes
2.

Systematic literature review

Coding protocol Age group study type of statistical test p value effect size (if possible) mediation/post-hoc analyses?

Results

Statistical power: Simulations

Statistical power is the probability that a particular statistical test will reject the null hypothesis.

I show simulations here that demonstrate issues with negative control groups, the use of post-hoc mediator variables

Basic power analysis

The key question of course for basic power analysis is the effect size that you would like to detect. If effects are large, then small samples are more justified. If effects are smaller, then small samples inflate the risk of false positives

Negative control groups

A common practice in infancy research is the addition of negative control experiments—that is, control groups that test an alternative explanation and for which the prediction is a failure. It is not unknown to see these control groups presented as distinct studies, with independent statistical tests performed. Then, a positive result in the primary experiment of interest and a failure to find a statistically significant result are used together to argue for the importance of the factor that was manipulated in the first experiment.

Of course, this experimental logic is not sound. These two experiments are in fact *different conditions* of the *same experiment*, because there is a manipulation whose causal importance is being assessed by the contrast in effects between groups. And since these two conditions are being compared, the difference between them must be tested statistically. The alternative strategy of (?, ?, ?)

Mediation and other post-hoc analyses

Consider a researcher who runs a standard infancy study with N=16 or N=24 and then goes on to explore a gender effect within that dataset. Even if

Figure 2 shows the results of a power analysis that assumes that the gender effect is equal in magnitude to the original effect (essentially, that one gender showed the predicted

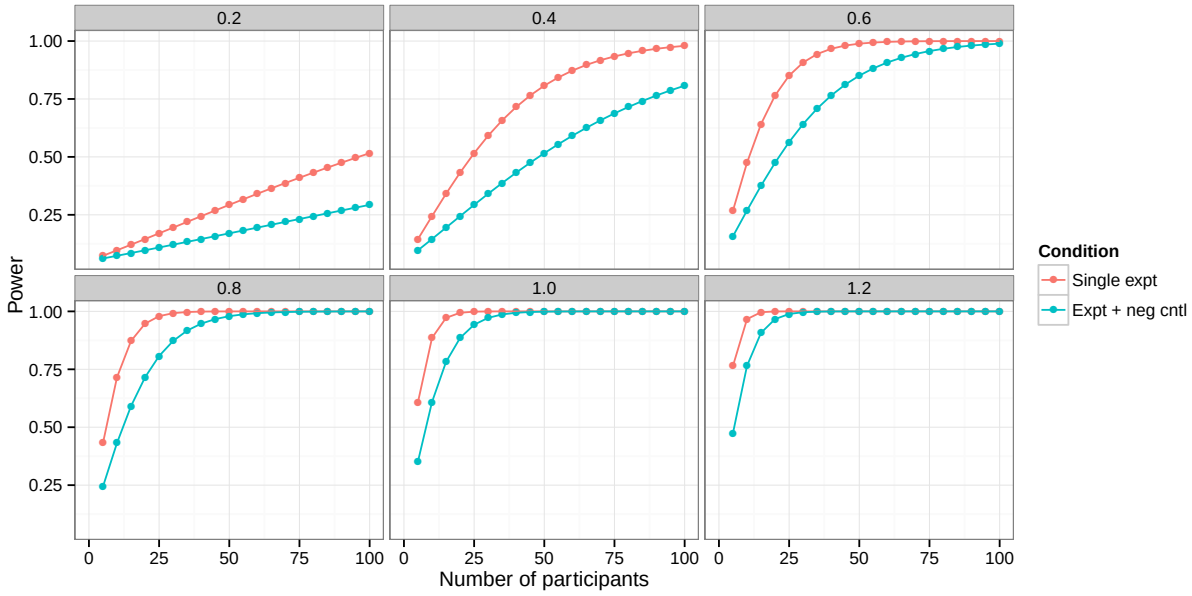

Figure 1. ....

effect and the other showed zero effect).

Recommendations

We close with a number of recommendations These are: 1) larger samples, 2) internal replications (especially including developmental replicates), 3) use of more transparent visualizations, 4) avoidance of post-hoc analysis, and 5) use of multi-trial paradigms.

Internal replications The strongest test of the reproducibility of a finding is a direct replication. Authors should be encouraged to perform and report such replications in their manuscripts, and reviewers should feel empowered to request them. Authors

One

Developmental replicates

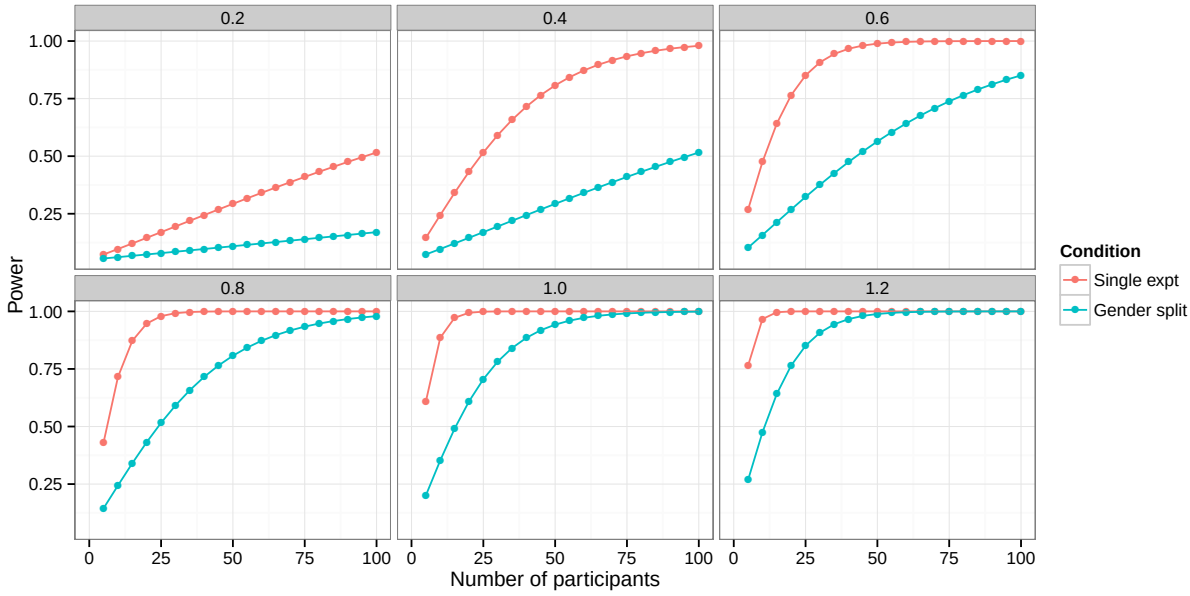

Figure 2. ....

Larger samples As is evidence from the preceding discussion, sample sizes in

Power analysis can be a helpful guide

In sum, the number of participants for a study will likely vary widely from study to study, but it will rarely be 16.

Visualization The most conventional visualization used in infancy research is the standard "dynamite" bar plot, with error bars representing the standard error of the mean. While they are clear and easy to read, these plots are relatively uninformative from the perspective of revealing features of the underlying dataset, as has been noted extensively in other literatures (?, ?). Authors should strongly consider moving towards plotting the individual participants'

In addition, the use of the standard error of the mean (SEM) is potentially misleading as it does not provide a good guide for inference. First, researchers should

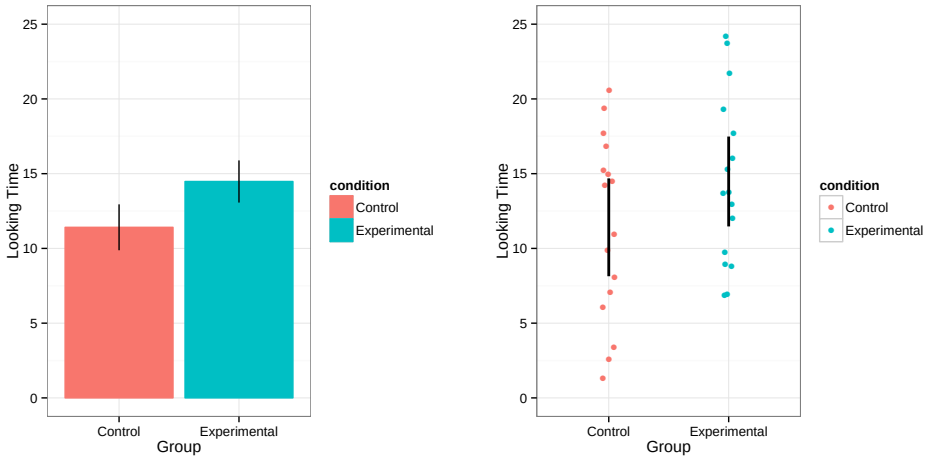

Figure 3. ....

move from SEM to 95% confidence intervals (?, ?). Second, researchers are often taught that non-overlapping standard errors are a meaningful indicator of the outcome of a statistical test—this is in fact false (?, ?). Such rules of thumb are doubly misleading when within-subject means are plotted side-by-side, since the appropriate statistical test is paired—hence SEM and CI are not inferentially relevant.

Post-hoc analyses Although such analyses are interesting, they are inevitably underpowered

Multi-trial paradigms

Conclusions: Implications for a literature